



Network Management

Network Management

**DE HOGESCHOOL
MET HET NETWERK**

Hogeschool PXL – Dep. PXL-IT – Elfde-Liniestraat 26 – B-3500 Hasselt
www.pxl.be - www.pxl.be/facebook



Intro

- configuratie nic:
 - via GUI: is mogelijk op een desktop
 - kan problemen geven als simultaan in CLI wordt geconfigureerd
 - via CLI: varieert afhankelijk van distributie
 - we bekijken Ubuntu



netplan

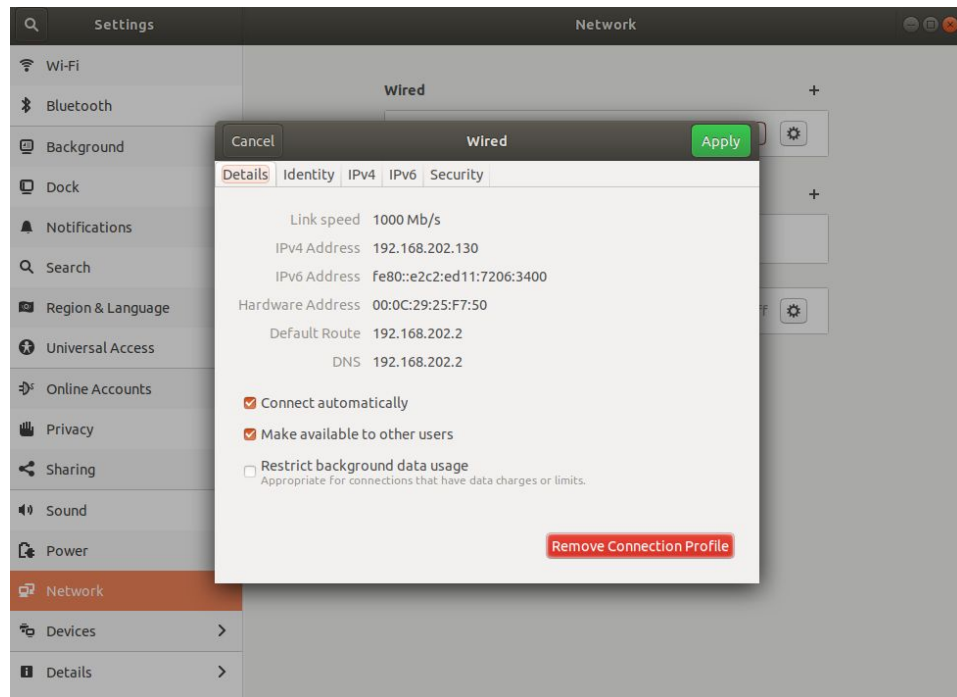
- vervangt het vroegere ifup/ifdown
 - Config in yaml file, afhankelijk van installatie:

Install Type	Renderer	File
Server ISO	systemd-networkd	/etc/netplan/01-netcfg.yaml
Cloud Image	systemd-networkd	/etc/netplan/50-cloud-init.yaml
Desktop ISO	NetworkManager	/etc/netplan/01-network-manager-all.yaml

```
student@ubuntu01:~$ cd /etc/netplan/  
student@ubuntu01:/etc/netplan$ ls  
01-network-manager-all.yaml  
student@ubuntu01:/etc/netplan$ cat 01-network-manager-all.yaml  
# Let NetworkManager manage all devices on this system  
network:  
  version: 2  
  renderer: NetworkManager  
student@ubuntu01:/etc/netplan$
```



Op de desktop



netplan

- vervangt het vroegere ifup/ifdown
 - Config in yaml file, afhankelijk van installatie:

Install Type	Renderer	File
Server ISO	systemd-networkd	/etc/netplan/01-netcfg.yaml
Cloud Image	systemd-networkd	/etc/netplan/50-cloud-init.yaml
Desktop ISO	NetworkManager	/etc/netplan/01-network-manager-all.yaml

1 netwerk kaart:

Renderer : NetworkManager

→ standaard op Desktop

→ voor wifi enz. via GUI

Ethernet (not wifi/bridge)

Naam: ens33 (later meer)

Geen vast ip-adres

dhcp4: ip-adres wordt via dhcp
geleverd over ipv4



Op de server

```
student@ubuntu-server01:~$ cd /etc/netplan/
student@ubuntu-server01:/etc/netplan$ ls
50-cloud-init.yaml
student@ubuntu-server01:/etc/netplan$ cat 50-cloud-init.yaml
# This file is generated from information provided by
# the datasource. Changes to it will not persist across an instance.
# To disable cloud-init's network configuration capabilities, write a file
# /etc/cloud/cloud.cfg.d/99-disable-network-config.cfg with the following:
# network: {config: disabled}
network:
  ethernets:
    ens33:
      addresses: []
      dhcp4: true
  version: 2
student@ubuntu-server01:/etc/netplan$
```

netplan

Beter voor een server: fixed ip (static)

gateway: alle trafic voor buiten het netwerk wordt naar de gateway (router) gestuurd

1 netwerk kaart:

Renderer: networkd

→ geen NetworkManager want
geen GUI

Ethernet (not wifi/bridge)

Naam: ens33 (later meer)

vaste instellingen voor
ip-adres met prefix
gateway

nameservers

search domain

dhcp4: geen dhcp

```
student@ubuntu01:~$ cat 50-cloud-init.yaml
# This file is generated from information provided by
# the datasource. Changes to it will not persist across an instance.
# To disable cloud-init's network configuration capabilities, write a file
# /etc/cloud/cloud.cfg.d/99-disable-network-config.cfg with the following:
# network: {config: disabled}
network:
  version: 2
  renderer: networkd
  ethernets:
    ens33:
      addresses: [172.16.226.137/24]
      gateway4: 172.16.226.2
      nameservers:
        addresses: [172.16.226.2, 8.8.8.8, 8.8.4.4]
        search: [sysnet.lan]
        # dhcp4: no
        # dhcp6: no
student@ubuntu01:~$
```



netplan

Wijzigingen in `/etc/netplan/*.yaml` worden doorgevoerd met
`sudo netplan apply`

Dit maakt ook de volgende file aan

`/run/systemd/network/<nr>-netplan-<nic>.network`

DHCP

```
student@ubuntu01:~$ cat /run/systemd/network/10-netplan-ens33.network
[Match]
Name=ens33

[Network]
DHCP=ipv4

[DHCP]
UseMTU=true
RouteMetric=100
student@ubuntu01:~$
```

STATIC

```
student@ubuntu01:~$ cat /run/systemd/network/10-netplan-ens33.network
[Match]
Name=ens33

[Network]
Address=172.16.226.137/24
Gateway=172.16.226.2
DNS=172.16.226.2
DNS=8.8.8.8
DNS=8.8.4.4
Domains=sysnet.lan
student@ubuntu01:~$
```

netplan

Wijzigingen in `/etc/netplan/*.yaml` worden doorgevoerd met
`sudo netplan apply`

Dit maakt ook de nodige instellingen in volgende file

`/etc/resolv.conf` → hier zie je enkel nog de DNS-search

→ resolving wordt nu gedaan door `systemd-resolve`

→ deze instellingen vind je nu terug in

`/run/systemd/resolve/resolv.conf`

OF via `systemd-resolve --status`



```
student@ubuntu-server01:~$ systemd-resolve --status | tail -4
DNS Servers: 172.16.226.2
              8.8.8.8
              8.8.4.4
DNS Domain: sysnet.lan
student@ubuntu-server01:~$
```

network status bekijken

Je kan de huidige instellingen ook bekijken met networkctl

Op de server

```
student@ubuntu-server01:~$ networkctl
```

IDX	LINK	TYPE	OPERATIONAL	SETUP
1	lo	loopback	carrier	unmanaged
2	ens33	ether	routable	configured

```
2 links listed.
```

```
student@ubuntu-server01:~$
```

```
student@ubuntu-server01:~$ networkctl status ens33
```

```
• 2: ens33
```

```
Link File: /lib/systemd/network/99-default.link
```

```
Network File: /run/systemd/network/10-netplan-ens33.network
```

```
Type: ether
```

```
State: routable (configured)
```

```
Path: pci-0000:02:01.0
```

```
Driver: e1000
```

```
Vendor: Intel Corporation
```

```
Model: 82545EM Gigabit Ethernet Controller (Copper) (PRO/1000)
```

```
HW Address: 00:0c:29:31:8f:3d (VMware, Inc.)
```

```
Address: 172.16.226.137
```

```
fe80::20c:29ff:fe31:8f3d
```

```
Gateway: 172.16.226.2 (VMware, Inc.)
```

```
DNS: 172.16.226.2
```

```
Search Domains: localdomain
```

```
student@ubuntu-server01:~$
```



Benaming netwerkkaarten

```
student@ubuntuserver01:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
p default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel stat
e UP group default qlen 1000
    link/ether 00:0c:29:9d:eb:4e brd ff:ff:ff:ff:ff:ff
    inet 172.16.226.137/24 brd 172.16.226.255 scope global dynamic ens33
        valid_lft 1030sec preferred_lft 1030sec
    inet6 fe80::20c:29ff:fe9d:eb4e/64 scope link
        valid_lft forever preferred_lft forever
student@ubuntuserver01:~$
```

In de laatste versies van ubuntu zijn de netwerkkaarten niet meer eth0, eth1, eth2, enz. genoemd

ens33: is hier de enige netwerkkaart.

→ Zou vroeger dus eth0 geweest zijn



Benaming netwerkkaarten

Vroeger werden de netwerkkaarten genummerd, in de volgorde dat de kernel deze zag tijdens het booten. Daarom kon het soms zijn dat eth0 en eth1 ineens gewisseld waren!

Nu wordt iedere netwerkkaart tijdens het booten hernoemd naar een eenduidige naam, die steeds hetzelfde is.

en -- ethernet sl -- serial line IP (slip) wl -- wlan ww -- wwan

b<number> -- BCMA bus core number ccw<name> -- CCW bus group name

o<index>[d<dev_port>] -- on-board device index number

s<slot>[f<function>][d<dev_port>] -- hotplug slot index number

x<MAC> -- MAC address

```
$ dmesg | grep -i eth
```

```
[ 3.050064] e1000 0000:02:01.0 eth0: (PCI:66MHz:32-bit) 00:0c:29:05:a3:e2
[ 3.050074] e1000 0000:02:01.0 eth0: Intel(R) PRO/1000 Network Connection
[ 3.057410] e1000 0000:02:01.0 ens33: renamed from eth0
```



Benaming netwerkkaarten

Het hernoemen van de netwerkkaarten naar een naam die steeds hetzelfde blijft is een goed idee. Indien je echter terug wilt naar de oude namen eth0, eth1, enzovoorts, kan je dat op volgende wijze:

```
$ sudo nano /etc/default/grub
```

```
GRUB_CMDLINE_LINUX="net.ifnames=0 biosdevname=0"
```

```
$ sudo grub-mkconfig -o /boot/grub/grub.cfg
```

```
$ sudo nano /etc/network/interfaces
```

```
auto eth0  
iface eth0 inet static  
    address 192.168.12.12
```

```
$ sudo reboot
```



ifconfig vs ip

- `ifconfig` is deprecated
 - `ip` is het vervangende commando

```
student@ubuntuserver01:~$ ifconfig ens33
ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.16.226.137 netmask 255.255.255.0 broadcast 172.16.226
    .255
    inet6 fe80::20c:29ff:fe9d:eb4e prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:9d:eb:4e txqueuelen 1000 (Ethernet)
    RX packets 315 bytes 354570 (354.5 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 198 bytes 14947 (14.9 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

student@ubuntuserver01:~$ ip a s ens33
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel stat
e UP group default qlen 1000
    link/ether 00:0c:29:9d:eb:4e brd ff:ff:ff:ff:ff:ff
    inet 172.16.226.137/24 brd 172.16.226.255 scope global dynamic ens33
        valid_lft 1782sec preferred_lft 1782sec
    inet6 fe80::20c:29ff:fe9d:eb4e/64 scope link
        valid_lft forever preferred_lft forever
student@ubuntuserver01:~$
```



ifconfig - deprecated

- informatie opvragen en wijzigingen aanbrengen
 - informatie opvragen zonder arguments: **alle nic's**

```
student@ubuntuServer:~$ ifconfig
ens33  Link encap:Ethernet  HWaddr 00:0c:29:52:5a:5b
        inet addr:172.16.110.128  Bcast:172.16.110.255  Mask:255.255.255.0
        inet6 addr: fe80::20c:29ff:fe52:5a5b/64 Scope:Link
        UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
        RX packets:56 errors:0 dropped:0 overruns:0 frame:0
        TX packets:53 errors:0 dropped:0 overruns:0 carrier:0
        collisions:0 txqueuelen:1000
        RX bytes:12834 (12.8 KB)  TX bytes:5903 (5.9 KB)

lo      Link encap:Local Loopback
        inet addr:127.0.0.1  Mask:255.0.0.0
        inet6 addr: ::1/128 Scope:Host
        UP LOOPBACK RUNNING  MTU:65536  Metric:1
        RX packets:160 errors:0 dropped:0 overruns:0 frame:0
        TX packets:160 errors:0 dropped:0 overruns:0 carrier:0
        collisions:0 txqueuelen:1
        RX bytes:11840 (11.8 KB)  TX bytes:11840 (11.8 KB)
```

Mac address

Subnet mask

IP address

Broadcast address



ip address

- informatie opvragen en wijzigingen aanbrengen
 - informatie opvragen zonder arguments: **alle nic's**

```
student@ubuntuServer:~$ ip address
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 00:0c:29:52:5a:5b brd ff:ff:ff:ff:ff:ff
    inet 172.16.110.128/24 brd 172.16.110.255 scope global ens33
        valid_lft forever preferred_lft forever
    inet6 fe80::20c:29ff:fe52:5a5b/64 scope link
        valid_lft forever preferred_lft forever
```

Mac address

Link state

Broadcast address

IP address

prefix



Of ip addr

of ip address show

of ip a

of ip a s

ip address show <nic>

- informatie opvragen en wijzigingen aanbrengen
 - informatie opvragen met arguments: 1 specifieke nic

```
student@ubuntuServer:~$ ip a s dev ens33
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 00:0c:29:52:5a:5b brd ff:ff:ff:ff:ff:ff
    inet 172.16.110.128/24 brd 172.16.110.255 scope global ens33
        valid_lft forever preferred_lft forever
    inet6 fe80::20c:29ff:fe52:5a5b/64 scope link
        valid_lft forever preferred_lft forever
```

`ip a s ens33` of `ip a s dev ens33`



ip addr show up

```
student@ubuntuServer:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens33: <BROADCAST,MULTICAST> mtu 1500 qdisc pfifo_fast state DOWN group default qlen 1000
    link/ether 00:0c:29:52:5a:5b brd ff:ff:ff:ff:ff:ff
    inet 172.16.110.128/24 brd 172.16.110.255 scope global ens33
        valid_lft forever preferred_lft forever
```

```
student@ubuntuServer:~$ ip a s up
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
```



'ip a s up' toont enkel de actieve linken

('down' bestaat niet als argument)

ip link up en ip link down

disablen van een nic

```
ip link set ens33 down
```

enablen van een nic

```
ip link set ens33 up
```

→ Dit leest de netplan-yaml file opnieuw uit



herstarten van de networkservice

```
sudo netplan apply
```

OF

```
sudo systemctl restart NetworkManager      (op Desktop)
```

```
sudo systemctl restart systemd-networkd     (op Server)
```

configuratie uit yaml-file in /etc/netplan wordt opnieuw toegepast



setting up IP address

Tijdelijke wijziging

```
student@ubuntuServer:~$ ip a s ens33 | grep 172
    inet 172.16.110.100/24 brd 172.16.110.255 scope global ens33
student@ubuntuServer:~$ sudo ip addr del 172.16.110.100/24 dev ens33
student@ubuntuServer:~$ ip a s ens33 | grep 172 ——— geen output=geen ip
student@ubuntuServer:~$ sudo ip addr add 172.16.110.101/24 dev ens33
student@ubuntuServer:~$ ip a s ens33 | grep 172
    inet 172.16.110.101/24 scope global ens33
```

<code>ip addr del 172.16.110.100/24 dev ens 33</code>	om 1 IP van de nic te verwijderen
<code>ip addr flush dev ens33</code>	om alle IPs van de nic te verwijderen
<code>ip addr add 172.16.110.101/24 dev ens33</code>	om 1 IP van de nic toe te voegen

ip link up → leest de yaml-file opnieuw uit !



```
student@ubuntuServer:~$ sudo ip link set ens33 down && sudo ip link set ens33 up
student@ubuntuServer:~$ ip a s ens33 | grep 172
    inet 172.16.110.100/24 brd 172.16.110.255 scope global ens33
```

hostname

Tijdelijke wijziging

Nieuwe naam zichtbaar bij het starten van een nieuwe shell

```
student@ubuntuServer:~$ sudo hostname nieuwenaaam
student@ubuntuServer:~$ bash
student@nieuwenaaam:~$ cat /etc/hostname
ubuntuServer
```

/etc/hostname is niet aangepast, dus bij een reboot opnieuw oude naam



hostnames mogen bestaan uit 64 letters, cijfers, (punten en) koppeltekens, maar niet eindigen met een koppelteken.

hostnamectl

Permanente wijziging

```
student@ubuntuServer:~$ cat /etc/hostname
ubuntuServer
student@ubuntuServer:~$ sudo hostnamectl set-hostname nieuwenaam
sudo: unable to resolve host nieuwenaam
student@ubuntuServer:~$ cat /etc/hostname
nieuwenaam
student@ubuntuServer:~$ bash
student@nieuwenaam:~$
```

/etc/hostname is aangepast, dus bij het starten van een nieuwe shell en het herstarten van de PC in de toekomst blijft de nieuwe naam behouden.

ENKEL ALS: in de file `/etc/cloud/cloud.cfg => preserve_hostname: true`

Default:

```
student@ubuntuserver01:~$ cat /etc/cloud/cloud.cfg | grep preserve_hostname
preserve_hostname: false
```

Je kan natuurlijk ook gewoon de file `"/etc/hostname"` aanpassen met vi of nano.



/etc/hosts

Aanpassen van /etc/hosts voor name-resolving (voor sudo)

- sudo doet voor ieder commando een name-resolving voor de hostname
 - Daarom is het belangrijk dat je de file /etc/hosts ook aanpast
 - anders heb je een lange timeout voordat een sudo commando wordt uitgevoerd

```
student@ubuntuServer:~$ cat /etc/hostname
ubuntuServer
student@ubuntuServer:~$ cat /etc/hosts
127.0.0.1    localhost
127.0.1.1    ubuntuServer

# The following lines are desirable for IPv6 capable hosts
::1         localhost ip6-localhost ip6-loopback
ff02::1     ip6-allnodes
ff02::2     ip6-allrouters
```



arp - deprecated

Toon alle entries

```
student@ubuntu01:~$ arp -a
_gateway (192.168.202.2) at 00:50:56:ec:9d:b4 [ether] on ens33
? (192.168.202.130) at 00:0c:29:25:f7:50 [ether] on ens33
```

```
student@ubuntu01:~$ ip a s dev ens33
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP g
    link/ether 00:0c:29:a8:0c:fb brd ff:ff:ff:ff:ff:ff
    inet 192.168.202.10/24 brd 192.168.202.255 scope global ens33
```

```
student@ubuntu01:~$ ping 192.168.202.132
PING 192.168.202.132 (192.168.202.132) 56(84) bytes of data.
64 bytes from 192.168.202.132: icmp_seq=1 ttl=64 time=0.403 ms
```

```
student@ubuntu01:~$ arp -a
? (192.168.202.132) at 00:0c:29:76:2f:ad [ether] on ens33
_gateway (192.168.202.2) at 00:50:56:ec:9d:b4 [ether] on ens33
? (192.168.202.130) at 00:0c:29:25:f7:50 [ether] on ens33
```



arp - deprecated

```
student@ubuntu01:~$ arp -a
? (192.168.202.132) at 00:0c:29:76:2f:ad [ether] on ens33
_gateway (192.168.202.2) at 00:50:56:ec:9d:b4 [ether] on ens33
? (192.168.202.130) at 00:0c:29:25:f7:50 [ether] on ens33
```

Delete entry

```
student@ubuntu01:~$ sudo arp -d 192.168.202.132
student@ubuntu01:~$ arp -a
_gateway (192.168.202.2) at 00:50:56:ec:9d:b4 [ether] on ens33
? (192.168.202.130) at 00:0c:29:25:f7:50 [ether] on ens33
```

```
student@ubuntu01:~$ ping 192.168.202.133
PING 192.168.202.133 (192.168.202.133) 56(84) bytes of data.
From 192.168.202.10 icmp_seq=1 Destination Host Unreachable
```

ping naar onbestaand IP

```
student@ubuntu01:~$ arp -a
? (192.168.202.133) at <incomplete> on ens33
_gateway (192.168.202.2) at 00:50:56:ec:9d:b4 [ether] on ens33
? (192.168.202.130) at 00:0c:29:25:f7:50 [ether] on ens33
```



ip neighbor (arp-table)

```
IP-NEIGHBOUR(8)                                Linux                                IP-NEIGHBOUR(8)
```

NAME

`ip-neighbour` - neighbour/arp tables management.

SYNOPSIS

```
ip [ OPTIONS ] neigh { COMMAND | help }
```

```
ip neigh { add | del | change | replace } { ADDR [ lladdr LLADDR ] [  
      nud STATE ] | proxy ADDR } [ dev DEV ]
```

```
ip neigh { show | flush } [ proxy ] [ to PREFIX ] [ dev DEV ] [ nud  
      STATE ]
```

```
STATE := { permanent | noarp | stale | reachable | none | incomplete |  
      delay | probe | failed }
```

DESCRIPTION

The `ip neigh` command manipulates neighbour objects that establish bindings between protocol addresses and link layer addresses for hosts sharing the same link. Neighbour entries are organized into tables. The IPv4 neighbour table is also known by another name - the ARP table.



ip neighbor (of ip n)

Toon alle entries

```
student@ubuntuserver01:~$ ip n show
192.168.202.133 dev ens33 FAILED
192.168.202.2 dev ens33 lladdr 00:50:56:ec:9d:b4 STALE
192.168.202.130 dev ens33 lladdr 00:0c:29:25:f7:50 DELAY
```

192.168.202.133 ping host unreachable
192.168.202.130 Connectie via ssh

```
student@ubuntuserver02:~$ ip a s ens33
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group
    link/ether 00:0c:29:76:2f:ad brd ff:ff:ff:ff:ff:ff
    inet 192.168.202.132/24 brd 192.168.202.255 scope global dynamic ens33
```

```
student@ubuntuserver01:~$ ping 192.168.202.132
PING 192.168.202.132 (192.168.202.132) 56(84) bytes of data.
64 bytes from 192.168.202.132: icmp_seq=1 ttl=64 time=0.547 ms
```

```
student@ubuntuserver01:~$ ip n show
192.168.202.132 dev ens33 lladdr 00:0c:29:76:2f:ad STALE
192.168.202.133 dev ens33 FAILED
192.168.202.2 dev ens33 lladdr 00:50:56:ec:9d:b4 STALE
192.168.202.130 dev ens33 lladdr 00:0c:29:25:f7:50 DELAY
```

Delete ent

```
student@ubuntuserver01:~$ sudo ip n del 192.168.202.132 dev ens33
student@ubuntuserver01:~$ ip n show
192.168.202.133 dev ens33 FAILED
192.168.202.2 dev ens33 lladdr 00:50:56:ec:9d:b4 DELAY
192.168.202.130 dev ens33 lladdr 00:0c:29:25:f7:50 REACHABLE
```



route - deprecated

```
student@ubuntuserver01:~$ route
Kernel IP routing table
Destination      Gateway          Genmask          Flags Metric Ref    Use Iface
default          _gateway        0.0.0.0          UG    0      0      0 ens33
192.168.202.0    0.0.0.0         255.255.255.0    U      0      0      0 ens33
student@ubuntuserver01:~$ netstat -r
Kernel IP routing table
Destination      Gateway          Genmask          Flags  MSS  Window  irtt Iface
default          _gateway        0.0.0.0          UG      0    0        0 ens33
192.168.202.0    0.0.0.0         255.255.255.0    U      0    0        0 ens33
```

Tijdelijk wijzigingen routering: (vb. default gateway)

```
sudo route add default gw 192.168.202.xx
```

```
(man route)
```



ip route

IP-ROUTE(8)

Linux

IP-ROUTE(8)

NAME

`ip-route` - routing table management

SYNOPSIS

`ip` [ip-OPTIONS] **route** { COMMAND | **help** }

`ip route` { **list** | **flush** } SELECTOR

`ip route get` ADDRESS [**from** ADDRESS **iif** STRING] [**oif** STRING] [**tos** TOS]

`ip route` { **add** | **del** | **change** | **append** | **replace** } ROUTE

SELECTOR := [**root** PREFIX] [**match** PREFIX] [**exact** PREFIX] [**table** TABLE_ID] [**proto** RTPROTO] [**type** TYPE] [**scope** SCOPE]

DESCRIPTION

`ip route` is used to manipulate entries in the kernel routing tables.

Route types:

unicast - the route entry describes real paths to the destinations covered by the route prefix



ip route (of ip r)

```
student@ubuntuserver01:~$ ip route  
default via 192.168.202.2 dev ens33 proto static  
192.168.202.0/24 dev ens33 proto kernel scope link src 192.168.202.10
```

Of `ip r show` of `ip r list`

Tijdelijk wijzigingen routing: (vb. default gateway)

```
sudo ip route add default via 192.168.14.xx
```



ping

Met ping wordt vaak de TCP/IP configuratie getest.

(ook traceroute, dig)

Ctrl-C

```
student@ubuntuserver01:~$ ping 192.168.202.132
PING 192.168.202.132 (192.168.202.132) 56(84) bytes of data.
64 bytes from 192.168.202.132: icmp_seq=1 ttl=64 time=0.711 ms
64 bytes from 192.168.202.132: icmp_seq=2 ttl=64 time=0.830 ms
64 bytes from 192.168.202.132: icmp_seq=3 ttl=64 time=0.871 ms
^C
--- 192.168.202.132 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2030ms
rtt min/avg/max/mdev = 0.711/0.804/0.871/0.067 ms
```



VMware: NAT-instelling in Workstation

The screenshot shows the VMware Workstation interface with the Virtual Network Editor window open for the VM 'OS_Ess_1J_UbServ1804'. The editor displays a table of network adapters and their configurations.

Name	Type	External Connection	Host Connection	DHCP	Subnet Address
VMnet1	Host-only	-	Connected	Enabled	192.168.152.0
VMnet5	Custom	-	-	-	192.168.220.0
VMnet8	NAT	NAT	Connected	Enabled	172.16.226.0

Below the table, the 'VMnet Information' section for VMnet8 is shown, with the 'NAT (shared host's IP address with VMs)' option selected. The 'Connect a host virtual adapter to this network' checkbox is checked, and the 'Host virtual adapter name' is 'VMware Network Adapter VMnet8'. The 'Use local DHCP service to distribute IP address to VMs' checkbox is also checked. The 'Subnet IP' is 172.16.226.0 and the 'Subnet mask' is 255.255.255.0.

At the bottom of the window, a warning message states: 'Administrator privileges are required to modify the network configuration.' with a 'Change Settings' button.

To direct input to this VM, click inside or press Ctrl+G.



VMware: NAT-instellingen op de host

interface vmnet8

NAT router en DNS Server

ip = 172.16.226.**2**

DHCP Server

ip = 172.16.226.**254**

```
PS C:\Users\gertv> Get-NetIPConfiguration *VMnet8

InterfaceAlias      : VMware Network Adapter VMnet8
InterfaceIndex      : 15
InterfaceDescription : VMware Virtual Ethernet Adapter for VMnet8
IPv4Address          : 172.16.226.1
IPv6DefaultGateway  :
IPv4DefaultGateway  :
DNSServer            : fec0:0:0:ffff::1
                    : fec0:0:0:ffff::2
                    : fec0:0:0:ffff::3

PS C:\Users\gertv>
```



VMware: NAT-instelling in de VM

Desktop: Een lease wordt aangevraagd door de NetworkManager bij het opstarten of herstarten van het netwerk van deze netwerkkaart.

```
student@ubuntudesktop01:~$ tail -16 /var/lib/NetworkManager/dhclient-e6c4d070-9d3e-3144-837c-92a3594e9379-ens33.lease
lease {
    interface "ens33";
    fixed-address 172.16.226.148;
    option subnet-mask 255.255.255.0;
    option routers 172.16.226.2;
    option dhcp-lease-time 1800;
    option dhcp-message-type 5;
    option domain-name-servers 172.16.226.2;
    option dhcp-server-identifier 172.16.226.254;
    option broadcast-address 172.16.226.255;
    option netbios-name-servers 172.16.226.2;
    option domain-name "localdomain";
    renew 3 2018/10/03 18:31:20;
    rebind 3 2018/10/03 18:43:52;
    expire 3 2018/10/03 18:47:37;
}
student@ubuntudesktop01:~$
```



VMware: NAT-instelling in de VM

Server: Een lease wordt aangevraagd door de networkd bij het opstarten of herstarten van het netwerk van deze netwerkkaart.

```
student@ubuntu-server01:~$ cat /var/run/systemd/netif/leases/2
# This is private data. Do not parse.
ADDRESS=172.16.226.137
NETMASK=255.255.255.0
ROUTER=172.16.226.2
SERVER_ADDRESS=172.16.226.254
NEXT_SERVER=172.16.226.254
T1=900
T2=1575
LIFETIME=1800
DNS=172.16.226.2
DOMAINNAME=localdomain
CLIENTID=ff2b9434c100020000ab11e422c0dcbce95e73
student@ubuntu-server01:~$
```

```
student@ubuntu-server01:~$ networkctl
IDX LINK                TYPE          OPERATIONAL SETUP
  1 lo                    loopback      carrier    unmanaged
  2 ens33                 ether         routable   configured

2 links listed.
student@ubuntu-server01:~$
```



VMware: NAT

Default GW

```
student@ubserv:~$ route
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
default 172.16.226.2 0.0.0.0 UG 0 0 0 ens33
172.16.226.0 * 255.255.255.0 U 0 0 0 ens33
```

```
student@ubserv:~$ dig www.pxl.be
```

```
; <<>> DiG 9.10.3-P4-Ubuntu <<>> www.pxl.be
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 18360
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; MBZ: 0005 , udp: 4000
;; QUESTION SECTION:
;www.pxl.be. IN A

;; ANSWER SECTION:
www.pxl.be. 5 IN A 192.168.4.20

;; Query time: 4 msec
;; SERVER: 172.16.226.2#53(172.16.226.2)
;; WHEN: Fri Oct 06 15:30:13 CEST 2017
;; MSG SIZE rcvd: 55
```

DNS

DNS - poort 53



ifconfig vs ip

COMPARING NET-TOOLS VS. IPRROUTE PACKAGE COMMANDS

NET-TOOLS COMMANDS

IPROUTE COMMANDS

arp -a

ip neigh

arp -v

ip -s neigh

arp -s 192.168.1.1 1:2:3:4:5:6

ip neigh add 192.168.1.1 lladdr 1:2:3:4:5:6 dev eth1

arp -i eth1 -d 192.168.1.1

ip neigh del 192.168.1.1 dev eth1

ifconfig -a

ip addr

ifconfig eth0 down

ip link set eth0 down

ifconfig eth0 up

ip link set eth0 up

ifconfig eth0 192.168.1.1

ip addr add 192.168.1.1/24 dev eth0

ifconfig eth0 netmask 255.255.255.0

ip addr add 192.168.1.1/24 dev eth0

ifconfig eth0 mtu 9000

ip link set eth0 mtu 9000

ifconfig eth0:0 192.168.1.2

ip addr add 192.168.1.2/24 dev eth0

netstat

ss

netstat -neopa

ss -neopa

netstat -g

ip maddr

route

ip route

route add -net 192.168.1.0 netmask 255.255.255.0 dev eth0

ip route add 192.168.1.0/24 dev eth0

route add default gw 192.168.1.1

ip route add default via 192.168.1.1

