

# AI & Robotics

State space and game AI

# Goals



## The **junior-colleague**

- can describe & explain in own words the position game AI as a subfield
- can describe in own words the link between tree search and game AI
- can explain in own words the differences between simple and more complex board games in context of game AI using real world games
- can explain in own words the term “contingency” of a problem
- can describe in own words how a game AI using tree search on an abstract level
- can explain in own words Minimax in context of game AI
- can explain in own words the time and space complexity of Minimax
- can implement Minimax for a given problem
- can explain in own words an improvement of Minimax
- can implement an improvement of Minimax
- can explain in own words Alpha-Beta pruning in context of game AI
- can explain in own words the best case and worst case gain of Alpha-Beta pruning in context of game AI
- can describe in own words the term Heuristic continuation in context of game AI and what problem it solves in context of Alpha-Beta pruning using a real world example
- can implement Alpha-Beta pruning for a given problem
- can implement Heuristic continuation for a given problem

# Why?

- One of the oldest subfields of AI
- Abstract and pure form of competition that seems to require intelligence
- Game playing is a special case of a search problem, with some new requirements.

# How?

- Simple board games
  - Easy to represent the states and actions
  - Very little world knowledge required!
  - “Contingency” problem:  
=> We do not know the opponents move!
  - The size of the search space:
    - Chess : +/- 15 moves possible per state, 80 plays =>  $15^{80}$  nodes in tree
    - Go : +/- 200 moves per state, 300 plays =>  $200^{300}$  nodes in tree
- More complex games
  - State space representation becomes increasingly difficult  
=> How to represent a game world?


# How?

## **Game playing algorithms:**

- Search tree only up to some depth bound
- Use an evaluation function at the depth bound
- Propagate the evaluation upwards in the tree


# MiniMax

- Consider a board game with 2 players:
  - **MAX** (AI player)
  - **MIN** (opponent)
- Each player alternates between taking a turn
- The game ends when either:
  - One of the 2 players reaches a winning state
  - No more moves are possible
- Assumptions:
  - Deterministic
  - Perfect information



# MiniMax

- Select a depth-bound (say: 2) and evaluation function
- Construct the tree up till the depth-bound
- Compute the evaluation function for the leaves




# MiniMax

- **MAX**-player wants to maximize **MIN**-player's ultimate score
- **MIN**-player wants to minimize **MAX**-player's ultimate score

=> Propagate the evaluation function Upwards:

- Take minima in **MIN**
- Take maxima in **MAX**





# MiniMax

```
init depthBound
function miniMax(board, depth):
    if depth == depthBound
        return eval(board)
    else if maximizer(depth)
        for each child c of board
            value = max(value, miniMax(child, depth + 1))
        return value
    else [minimizer]
        for each child c of board
            value = min(value, miniMax(child, depth + 1))
        return value
```

# MiniMax

## Analysis

- Time complexity

=> Same as iterative deepening (search bounded by depth  $m$ ):  **$O(b^m)$**

- Space complexity


=> Same as iterative deepening (search bounded by depth  $m$ ):  **$O(b*m)$**

# MiniMax: Tic-Tac-Toe


- 2 players: X and O
- State representation of the board: i.e. matrix
- Production rules:
  - X move: place X on a free spot on the board
  - O move: place O on a free spot on the board
- Start state: empty board
- Goal state:
  - 3 X's in a row
  - 3 O's in a row
  - Full board

# MiniMax: Tic-Tac-Toe

- **MAX**-player (X) wins: +10
- **MIN**-player (O) wins: -10




```
function eval(board, depth):  
    if maximizer(depth)  
        return 10  
    else if  
        minimizer(depth)  
        return -10  
    else  
        return 0
```



# MiniMax: Tic-Tac-Toe


## MAX-player X's turn in State 1

- State 1: generates states 2, 3, and 4 and calls minimax on those states
- State 2: goal state  
=> X win: return +10 to state 1
- State 3: generates states 5 and 6 and calls minimax on them
- State 4: generates states 7 and 8 and calls minimax on them
- State 5: goal state  
=> O win: return -10 to state 3
- State 7: goal state  
=> O win: return -10 to state 4
- State 6: generates state 9 and calls minimax on it
- State 9: goal state  
=> X win: return +10 to state 3
- State 8: generates state 10 and calls minimax on it
- State 10: goal state  
=> X win: return +10 to state 4



# MiniMax: Tic-Tac-Toe

- State 6 and 8: generate states 9 and 10 and call minimax on them
- State 9 and 10: goal states  
=> return +10 to states 6 and 8  
=> return +10 to states 3 and 4
- State 3 and 4: O's turn  
=> **MIN**imize score: **MIN**(-10, +10) = -10  
=> states 3 and 4 return -10
- State 1: X's turn  
=> **MAX**imize score **MAX**(+10, -10, -10) = +10  
=> Choose State 2




# MiniMax: Improvement

## Problem

Early demise: algorithm doesn't differentiate between an early and a late defeat

=> O player could choose state 2, 4 or 5 instead of state 3




# MiniMax: Improvement

## Solution

Delay demise: take depth into account for evaluation score

```
function eval(board, depth):
    if maximizer(depth)
        return 10 - depth
    else if minimizer(depth)
        return depth - 10
    else
        return 0
```






# Alpha-Beta pruning

- Optimization for MiniMax
- Instead of:
  - first creating the entire tree (up to depth-level)
  - then doing all propagation
- Interleave the generation of the tree and the propagation of values.

=> some of the obtained values in the tree will provide information that other (non-generated) parts are redundant and do not need to be generated.


# Alpha-Beta pruning

- Generate the tree depth-first, left to right
- Propagate final values of nodes as initial estimates for their parent node
- The MIN-value (1) is already smaller than the MAX-value of the parent (2)
- The MIN-value can only decrease further
- The MAX-value is only allowed to increase
- No point in computing further below this node



# Alpha-Beta pruning


- The values at **MAX** nodes are **Alpha-values**
- The values at **MIN** nodes are **Beta-values**



# Alpha-Beta pruning

If an **Alpha-value** is larger or equal to the **Beta-value** of a descendant node:


=> Stop generation of the children of the descendant



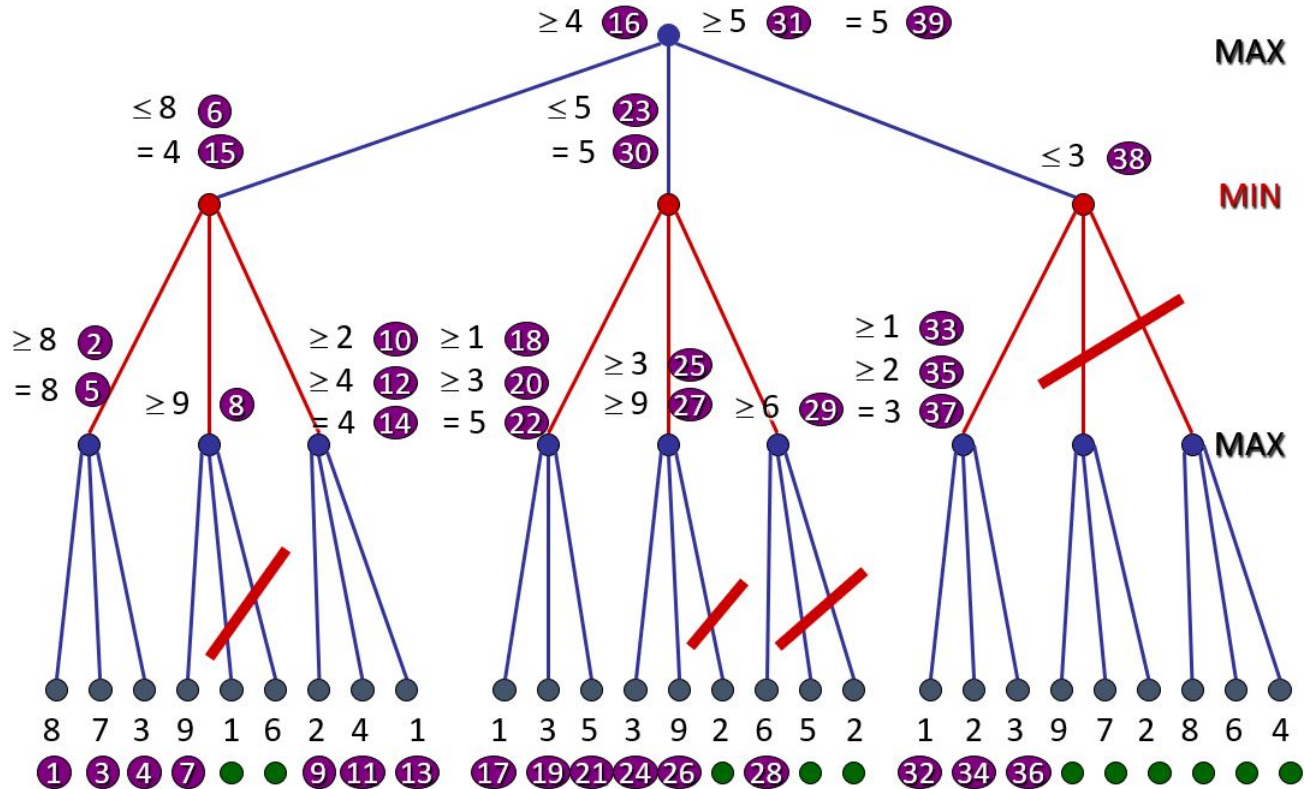
# Alpha-Beta pruning

If an **Beta-value** is smaller or equal to the **Alpha-value** of a descendant node:

=> Stop generation of the children of the descendant



# Alpha-Beta pruning



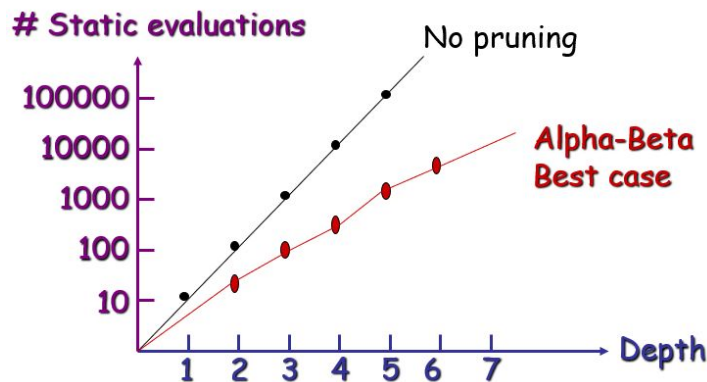

# Alpha-Beta pruning

## Best case gain:

- At every layer, the best node is the **left-most one**:  
=> only ■ is explored
- Evaluations saved:  
 $O(b^{d/2})$

## Worst case gain:


- No improvement



# Alpha-Beta pruning

**Best case gain:**


Example of a perfectly ordered tree





# Alpha-Beta pruning: Problem

- Depth-bound is limiting factor
  - It's preferable to delay disasters, but they are not prevented
- => possible solution:  
heuristic continuation



# Alpha-Beta pruning: Heuristic continuation

- Change behaviour in certain situations
    - Strategically crucial:
      - e.g. chess: king in danger, pawn can convert to queen, etc
- => extend search beyond the depth bound

