

# AI & Robotics

How good is my model?

# Goals



## The **junior-colleague**

- can explain in their own words the difficulty of describing the “goodness” of a ML model
- can explain in their own words the metric used for classification
- can explain in their own words the metrics used for regression
- can describe how to go about measuring the “goodness” of an ML model
- can explain overfitting and underfitting in their own words
- can explain the tradeoff between bias and variance
- can explain the importance of dividing your data sets in training, validation and test sets
- can explain the purpose of the training set
- can explain the purpose of the validation set
- can explain the purpose of the test set

# What does it mean for a model to be good?

- How do we measure “goodness”?
- How do we analyze it?
- How can we improve it?

# How do we measure “goodness”?

## **Problem situation: a unique ML problem**


- No indication what a good model looks like
- The goodness of a model is relative
- The best that we can do is to compare the performance of ML models on your specific data to other models also trained on the same data

# How do we measure “goodness”?

A perfect model is practically impossible!

All predictive modeling problems have prediction error, coming from:

- Missing values
- Noise in the data
- Stochastic nature of Machine Learning



Even robots can be wrong. :\*(-(

Do not trust AI blindly.

# Metrics: Classification


## Accuracy

	Actual positive	Actual negative
Predicted positive	True Positive (TP)	False Positive (FP)
Predicted negative	False Negative (FN)	True Negative (TN)

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

=> 100% accuracy is the best!

# Metrics: Regression



Measure the distance between the predicted value and the actual value that the model was trained to predict

**(log/root)MSE:** the average amount by which the model will be off

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

=> 0.0 error is the best!

# Metrics: Regression

Measure the distance between the predicted value and the actual value that the model was trained to predict

**R<sup>2</sup>**: the proportion of the variance in the dependent variable that is predictable from the independent variable(s)


- E.g R<sup>2</sup> of a model is 1.0 => all movements of the dependent variable are completely explained by movements of the input
- E.g. R<sup>2</sup> of a model is 0.50 => half of the observed variation can be explained by the model's inputs

=> 1.0 R<sup>2</sup> score is the best!


$$R^2 = \frac{\text{Explained Variation}}{\text{Total Variation}}$$



# Metrics: Regression



Explained variation



Alternatively: 
$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$

# But wait, there's more!

There are even more metrics for different problems (e.g. recommendation systems):


- Precision
- Recall
- ROC analysis
- F1 score
- Etc.

=> will be covered later in this educational programme


# How do we measure “goodness”?

1. Define a baseline
  - Mean outcome value for a regression problem (the average value)
  - Mode outcome value for a classification problem (the value that appears most often)
2. Compare your trained model to the baseline


# Overfitting vs. underfitting



Underfitted




Good Fit/Robust




Overfitted


# Overfitting vs. underfitting



**overfitting**




**underfitting**




**Good balance**

# Overfitting vs. underfitting



# Overfitting vs. underfitting



# Overfitting vs. underfitting



Steve Maine

@smaine

TIL that changing random stuff until your program works is "hacky" and "bad coding practice" but if you do it fast enough it's "[#MachineLearning](#)" and pays 4x your current salary

6:40 PM · 10 May 18

629 Retweets 1,692 Likes





# Train, Validation, Test



# Train, Validation, Test

## **Training set**

- The dataset used to train the model
- The model learns from this data


## **Validation set**

- The dataset used to evaluate the model during training and optimization
- Used while fine tuning the model hyperparameters
- The model doesn't learn from this data

## **Test set**

- The dataset used to evaluate the model after training
- The model doesn't learn from this data

# Train, Validation, Test



- Training error  $\neq$  Validation error
- Use the validation set to tune the model

# How to split?



- Full set of samples: usually [80/20]::[train/test]
  - Training set: usually [80/20]::[train/validation]
  - It's important to choose good test and validation sets
    - Make sure there's a good spread in the test samples
    - Make sure your validation set conforms to the test set
- => more on this later

