

AI & Robotics

Data & Optimization



Goals



The **junior-colleague**

- can explain the link between data & optimization
- can describe the reasons for understanding your data
- can explain how to prepare and clean a dataset for training a Machine Learning model (categorical variables, dates, missing values)
- can explain one-hot encoding
- can explain the difference between parameters and hyperparameters
- can list and describe 3 ways to tune the hyperparameters of a Machine Learning model
- can list and describe 3 hyperparameters of a Random Forest
- can explain how to acquire the out-of-bag score for a Random Forest and why it's useful
- can explain why feature importance is important in the context of Machine Learning
- can explain the concept of Occam's Razor in the context of Machine Learning
- can describe ways to speed up the training process of a Random Forest

Intro



- Data & Optimization are inextricably linked
- Understanding your data gives you insights into optimization


Data




Data

... is the key ingredient for all Machine Learning problems!

The Cambrian Explosion...of Data



Traditional Data Sources Dominate, But Many New Sources Are Planned







Multiple responses allowed

© 2014 Gartner, Inc. and/or its affiliates. All rights reserved.

Processing data

Pandas


- Python library for data manipulation and analysis
- Highly optimized for performance
- Seamless integration with *numpy*



Understanding data

- Important to look at your data
 - Understand how it's stored
 - Understand the format and types of values

=> But! Don't get biased
- Important to get an idea of what your data is about
 - => find correlations
 - => look for ways to optimize your model
- Important to note what metric is being used for a project



Cleaning and preparing data

Categorical variables

- Categorical variables usually get stored as strings in pandas
 - => inefficient
 - => not numerical (impossible to process for most ML algorithms)
- Convert to “categories”
 - Will be treated as numerical data internally
 - Optional ordering

=> Make sure you use the same procedure for your different datasets, otherwise the categories won't match

Cleaning and preparing data

	A	B	C	D	E	F	G
1	Plot: 3						
2	Date collected	Species	Sex	Weight	Month	Day	Year
3	1/8	PF	M	7	=MONTH(A3)	=DAY(A3)	=YEAR(A3)
4	2/18	OT	M	24	2	18	2015
5	2/19	OT	F	23	2	19	2015
6	3/11	NA	M	232	3	11	2015
7	3/11	OT	F	22	3	11	2015
8	3/11	OT	M	26	3	11	2015
9	3/11	PF	M	8	3	11	2015
10	4/8	NA	F		4	8	2015
11	5/6				5	6	2015
12	5/18	NA	F	182	5	18	2015
13	6/9	OT	F	29	6	9	2015
14	7/8	NA	F	115	7	8	2015
15	7/8	NA	M	190	7	8	2015

“Dates”

- Contain lots of information
 - Day, Month, Year, DayOfWeek
 - Seasonal variation
 - Temporal continuity
- Separate data type in pandas

Cleaning and preparing data

Handling missing values

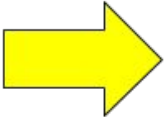
- Missing categorical values get a numeric code
- Missing continuous values get assigned the median (or other measure) for that feature
- Optionally create extra columns (features) for missing categorical data

=> Make sure you use the same procedure for your different datasets, otherwise the categories won't match

Cleaning and preparing data

One-hot encoding

- Necessary for some ML algorithms
- Random Forests don't require One-hot encoding



Color
Red
Red
Yellow
Green
Yellow


Red	Yellow	Green
1	0	0
1	0	0
0	1	0
0	0	1

Cleaning and preparing data

One-hot encoding


Random Forests don't require One-hot encoding

=> but, can be useful for new insights into your data



Optimization

Optimization: parameters vs. hyperparameters



- Parameters
 - Learned during training
- Hyperparameters
 - Set by the data scientist before training
 - **Validation set!!**

Optimization: hyperparameter tuning

How?

- Trial and error:
 - => Use metrics and graphs to estimate good hyperparameters
 - => Gives you good insights into the data
- Grid search
 - => Just try every possible combination of parameters and values
 - => Guaranteed to find the optimal hyperparameters
- Random search
 - => Try random parameters and values
 - => Not guaranteed to find the optimal hyperparameters

Depending on the size of the problem, tuning can require **A LOT** of computational resources

Optimization: Random Forests

- Scikit-learn hyperparameters
 - `n_estimators`
 - `max_depth`
 - `min_samples_split`
 - `min_samples_leaf`
 - `min_weight_fraction_leaf`
 - `max_features`
 - `max_leaf_nodes`
 - `min_impurity_decrease`
 - `min_impurity_split`
 - ...
- Check the documentation for details

Optimization: Random Forests

n_estimators

- Number of trees
- Incrementally increase the amount of trees
- Use graphs + R^2 to see at which point it stops helping
- Use fewer trees when experimenting

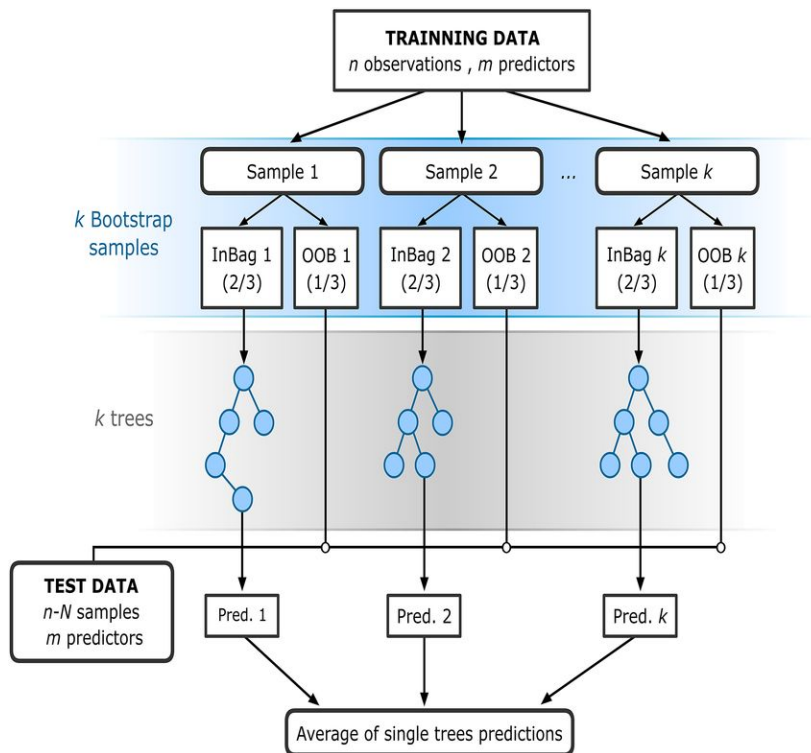
min_samples_leaf

- The minimum amount of samples in each leaf node
- Reduces the chance of overfitting

max_features


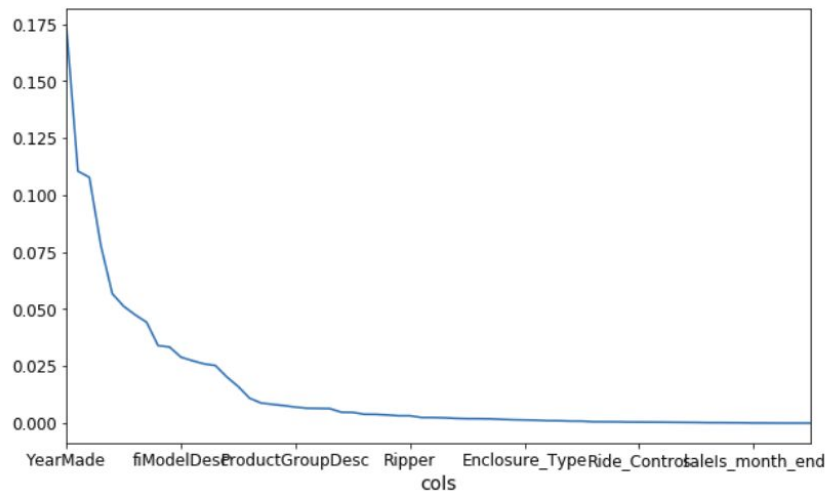
- The number of features to choose a splitpoint from at each node in a tree
- Increases randomness

Optimization: Random Forests - Out of bag



- What if our dataset is too small?
- For every tree: use the rows not used to train the model as a validation set
=> different validation set for each tree
- Average over all trees where that row was not used for training
- The more trees you have, the higher the chance of each sample appearing in valid set of multiple trees
- Less useful than a true validation set
- Is statistically significant
=> good overfitting check


Optimization: Feature Importance



Optimization: Feature Importance

- One of the most important optimization techniques
- Remove less important features, so the important ones stand out more, because there may have been collinearity
- How to calculate feature importance: Shuffle the rows of that feature, compare the R^2 of this model to the original model
- Remarks:
 - Shuffle method can be misleading if 2 attributes are closely related (correlated)
=> Checking for feature similarity can help
 - The model indicates Feature x is important, yet the client says that it doesn't make sense
=> Client may not understand his/her own data (data leakage)

Optimization: Feature similarity



1. Visualize potential correlations (i.e. hierarchical clustering graph)
2. Check how the model behaves when removing correlated features one at a time
3. Remove the features that have a negligible impact



How to measure similarity or correlation?

=> Correlation coefficient (R^2) / Spearman's R (Doesn't really matter what, as long as it measures correlation)

Optimization: Occam's Razor

Optimal tree

Overfit tree



Accuracy on training = 70%


Accuracy on test = 70%

Accuracy on training = 90%

Accuracy on test = 65%

- Simple models >>> complex models
- Similar principle in science: Occam's Razor
 - "Simpler solutions are more likely to be correct than complex ones"
 - "Select the solution with the fewest assumptions"
- Decision trees & Random Forests
 - => more complex trees have higher probability of overfitting the data set

Optimization: training speedup



Subsampling

Pick a different random subset of samples for each tree

=> given enough trees, the model can still see all the data

Optimization: closing remarks

- If something goes wrong in Machine Learning
 - No error message
 - Your model is just slightly less good
- Domain knowledge can help but not necessarily

